

Developments for ‘2048’ Project

This write-up documents the entire development process of a 2048 agent, including test results, proposed improvements, and design decisions.

MinMax Tree Search

1. Version / Baseline (`MinMax.py`) : MinMax with Simple Heuristic

- Method:
Minimax Tree Search; Simple Heuristic
- Results:

Average Score	Maximum Score	Median Score	2048 Percent	8192 Percent	Number of Games
9,872.67	25,876	7,578	1.4	0.0	72

- Observations:
 - MinMax is a poor match for 2048’s stochastic nature. It treats the random tile spawns too pessimistically, which limits average performance.
- Plan:
 - Replace Minmax Tree Search with ExpectiMax Tree Search

ExpectiMax Tree Search

2. Version (`MyAgent_Version02.py`) : ExpectiMax with Simple Heuristic

- Methodological improvements over the previous version:
Minmax Tree Search → ExpectiMax Tree Search
 - The “max” nodes are the AI choosing a move; the “chance” nodes represent random tile spawns.
 - At a max node, the action with the highest expected value is taken; at a chance node, the expectation (weighted average) over all random outcomes is computed.
- Results:

Average Score	Maximum Score	Median Score	2048 Percent	8192 Percent	Number of Games
10,317.12	26,460	10,632	3.2	0.0	250

- Observations:
 - Replacing MinMax with ExpectiMax gave a clear, measurable improvement: ExpectiMax models random tile spawns correctly and boosts the median and 2048 rate.
 - The gain is modest (~4.5%) because the heuristic still remains simple; the search model improved but leaf evaluation still limits quality.
 - This confirms the hypothesis that modeling chance is necessary but insufficient – better leaf evaluation is the main lever.
- Plan:
 - Replace the simple heuristic by an extended heuristic.

3. Version (`MyAgent_Version03.py`) : ExpectiMax with Extended Heuristic

- Methodological improvements over the previous version:
Simple Heuristic → Extended Heuristic
 - For the heuristic, a composite score is built by summing several heuristic components in a weighted linear manner.
 1. Heuristic component (`base_game_score`)
 - It is the game’s own score (sum of merged tile values that have been created so far).
 - It gives baseline progress toward winning.
 2. Heuristic component (`corner_score`)
 - It is a big positive bonus if the largest tile sits in one of the four corner indices; otherwise a penalty (and a special-case penalty if the board is empty).
 - The standard 2048 strategy is to keep the max tile in a stable corner to avoid breaking the ordered gradient and to reduce the chance of losing mobility.

- 3. Heuristic component (`empty_tile_score`)
 - It is the weighted number of empty tiles on the board. More empty cells mean more legal moves and higher chance to combine tiles later; empties strongly correlate with survival.
 - Results:

Average Score	Maximum Score	Median Score	2048 Percent	8192 Percent	Number of Games
24,398.63	59,232	26,020	55.0	0.0	131
 - Observations:
 - Introducing structured heuristic components (base score, corner, empties) caused a dramatic improvement. This shows heuristic design is the largest driver of performance so far.
 - The 2048 attainment rate jumped to 55% - strong evidence that encouraging corner-max and maintaining empties pays off.
 - Median and mean are both much higher and max also increased, indicating the heuristic improves both typical and best-case play.
 - Plan:
 - Replace the extended heuristic by a further-extended heuristic.
- 4. Version (`MyAgent.py`): ExpectiMax with Further-Extended Heuristic**
- Methodological improvements over the previous version:

Extended Heuristic → Further-Extended Heuristic

 - For the heuristic, a composite score is built by summing several heuristic components in a weighted linear manner.
 - Heuristic component (`base_game_score`): *Refer to version 3 for more details.*
 - Heuristic component (`corner_score`): *Refer to version 3 for more details.*
 - Heuristic component (`empty_tile_score`): *Refer to version 3 for more details.*
 - Heuristic component (`snake_pattern_score`)
 - Is a weighted sum of tile values arranged in a predefined snake ordering (weights in `SNAKE_WEIGHTS`), then normalized by `max_tile_exponent`.
 - The intent is to reward boards that follow a strong descending snake gradient (largest tile at the start of the snake).
 - The “snake” (or gradient) arrangement keeps tiles ordered and makes merges predictable while preserving the max tile location.
 - Heuristic component (`monotonicity_score`)
 - It counts rows/columns that are monotonic (non-decreasing or non-increasing) and multiplies by `MONOTONICITY_WEIGHT`.
 - Monotonic rows/columns indicate a gradient where tiles consistently increase or decrease across a line – that reduces the chance of blocking merges and helps keep the big tile in place.
 - Heuristic component (`merge_score`)
 - It counts adjacent equal non-zero tiles (horizontally and vertically) – each potential merge gets `MERGE_WEIGHT`.
 - This indicates immediate opportunities to increase score and create higher-value tiles. A board with many adjacent equal tiles is tactically stronger.
 - Heuristic component (`smoothness_score`):
 - This penalizes large differences between neighbouring tiles (summing absolute differences of exponents along rows and columns), multiplied by `SMOOTHNESS_SCALE`.
 - Smaller gradients are less likely to leave isolated high tiles that are hard to merge; smoothness encourages gradual value transitions.
 - Results:

Average Score	Maximum Score	Median Score	2048 Percent	8192 Percent	Number of Games
41,182.10	82,760	35,712	89.6	0.0	250
 - Observations:
 - Adding snake pattern, monotonicity, merge and smoothness components produced another substantial jump. These features capture long-term board geometry and merge likelihood.

- Very high 2048 rate ($\approx 90\%$) shows the agent reliably reaches the basic goal. Max score nearly doubled from Version 03, indicating better deep play.
- Median < mean and a very large max indicate there is still high variance / room for spectacular runs; but typical play (median) is now strong.
- Plan:
 - Add methodology for storing the best move.

5. Version (MyAgent_Version05.py) : ExpectiMax with Further-Extended Heuristic and Best Move Memory

- Methodological improvements over the previous version:
 - + Best Move Memory
 - Remembers the single move that the search previously judged best and tries it first next time.
 - If the search tries the most promising move first, it often finds a strong (or the strongest) line earlier. In MiniMax/Alpha-Beta this increases pruning; in ExpectiMax (here) it helps the search discover high values earlier which should improve the anytime result under a time limit.
- Results:

Average Score	Maximum Score	Median Score	2048 Percent	8192 Percent	Number of Games
39,845.09	75,488	35,782	77.3	0.0	77
- Observations:
 - Introducing a single “best move memory” produced a small decline in average and a noticeable drop in 2048 rate ($89.6 \rightarrow 77.3$).
 - One possible reason for the drop is stale-move bias: storing one global “best move” (rather than per-state/transposition) can bias the search toward moves that were good under different board contexts, reducing adaptability.
- Outlook:
 - During development, various versions of move ordering were tested, but none were found to improve the overall performance of the agent.
 - Hybrid Search: Run full ExpectiMax to some shallow depth d , then stop and evaluate the leaf with a learned value function (neural net) instead of the handcrafted heuristic. This should give more accurate leaf evaluations and enables to search deeper effectively.

Summary of Results

Ver-sion	Description	Avg. Score	Relative Change in ‘Avg. Score’	Max. Score	Med. Score	2048 %	8192 %	Num. of Games
1	Minimax Tree Search with Simple Heuristic	9,872.67	-	25,876	7,578	1.4	0.0	72
2	ExpectiMax Tree Search with Simple Heuristic	10,317.12	4.50%	26,460	10,632	3.2	0.0	250
3	ExpectiMax Tree Search with Extended Heuristic	24,398.63	136.49%	59,232	26,020	55.0	0.0	131
4	ExpectiMax Tree Search with Further-Extended Heuristic	41,182.10	68.79%	82,760	35,712	89.6	0.0	250
5	ExpectiMax Tree Search with Further-Extended Heuristic, Best Move Memory,	39,845.09	-3.25%	75,488	35,782	77.3	0.0	77

Discussion of Results

The heuristic design matters most. Replacing MinMax with ExpectiMax gave small gains (Version 01 \rightarrow Version 02); adding relevant heuristic features (Version 02 \rightarrow Version 03 and Version 03 \rightarrow Version 04) produced very large improvements. Structural features (snake, monotonicity, smoothness, merge potential) are highly effective. The search model matters too, but mainly insofar as it exploits a good leaf evaluator. ExpectiMax + good heuristic is far better than Minimax + simple heuristic. But even at Version 04 the distribution is skewed: high maxima indicate the agent can occasionally exploit excellent sequences. Median and mean analysis together shows improvements are both in typical games and best-case runs.